

# Homework 6

```
library(tidymodels)

## -- Attaching packages ----- tidymodels 0.2.0 --

## v broom      0.8.0    v recipes     0.2.0
## v dials      0.1.1    v rsample     0.1.1
## v dplyr      1.0.9    v tibble      3.1.7
## v ggplot2    3.3.6    v tidyverse   1.2.0
## v infer       1.0.0    v tune        0.2.0
## v modeldata   0.1.1    v workflows   0.2.6
## v parsnip     0.2.1    v workflowsets 0.2.1
## v purrr       0.3.4    v yardstick   0.0.9

## -- Conflicts ----- tidyverse_conflicts() --
## x purrr::discard() masks scales::discard()
## x dplyr::filter()  masks stats::filter()
## x dplyr::lag()    masks stats::lag()
## x recipes::step() masks stats::step()
## * Search for functions across packages at https://www.tidymodels.org/find/

library(ISLR)
library(ISLR2)

##
## Attaching package: 'ISLR2'

## The following objects are masked from 'package:ISLR':
## 
##     Auto, Credit

library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.1 --

## v readr      2.1.2    vforcats 0.5.1
## v stringr    1.4.0

## -- Conflicts ----- tidyverse_conflicts() --
## x readr::col_factor() masks scales::col_factor()
## x purrr::discard()   masks scales::discard()
## x dplyr::filter()    masks stats::filter()
## x stringr::fixed()   masks recipes::fixed()
## x dplyr::lag()       masks stats::lag()
## x readr::spec()      masks yardstick::spec()
```

```
library(glmnet)

## Loading required package: Matrix

##
## Attaching package: 'Matrix'

## The following objects are masked from 'package:tidyverse':
##       expand, pack, unpack

## Loaded glmnet 4.1-4

library(janitor)

##
## Attaching package: 'janitor'

## The following objects are masked from 'package:stats':
##       chisq.test, fisher.test

library(rpart.plot)

## Loading required package: rpart

##
## Attaching package: 'rpart'

## The following object is masked from 'package:dials':
##       prune

library(ranger)
library(corrplot)

## corrplot 0.92 loaded

library(vip)

##
## Attaching package: 'vip'

## The following object is masked from 'package:utils':
##       vi
```

```
library(xgboost)

## 
## Attaching package: 'xgboost'

## The following object is masked from 'package:dplyr':
## 
##     slice

tidymodels_prefer()

poke <- read.csv('Pokemon.csv')
pokemon <- clean_names(poke)

pokemon <- pokemon%>%filter(type_1 == 'Bug' | type_1== 'Fire' | type_1== 'Grass' | type_1== 'Normal' | 


pokemon$legendary <- as.factor(pokemon$legendary)

pokemon$type_1 <- as.factor(pokemon$type_1)
pokemon$generation <- as.factor(pokemon$generation)

set.seed(608)
pokeSplit <- initial_split(pokemon, prop = 0.80,
                             strata = type_1)
poke_train <- training(pokeSplit)
poke_test <- testing(pokeSplit)

poke_fold <- vfold_cv(poke_train, v = 5, strata = type_1)
pokeRecipe <- recipe(type_1 ~ legendary + generation + sp_atk + attack + speed + defense + hp + sp_def

poke_train %>% select(where(is.numeric)) %>% select(-x, -total) %>% cor() %>% corrplot(method = 'number
```



None of the relationships are very high, however, defense and hp, speed and defense, speed and hp, sp-def and speed, all have especially low correlations

```
tree_spec <- decision_tree() %>%
  set_engine("rpart")


class_tree_spec <- tree_spec %>%
  set_mode("classification")

class_tree_wf <- workflow() %>%
  add_model(class_tree_spec %>% set_args(cost_complexity = tune())) %>%
  add_recipe(pokeRecipe)

param_grid <- grid_regular(cost_complexity(range = c(-3, -1)), levels = 10)

tune_res <- tune_grid(
  class_tree_wf,
  resamples = poke_fold,
  grid = param_grid,
  metrics = metric_set(roc_auc)
)

autoplot(tune_res)
```



It performs better with a smaller cost-complexity parameter

```
collect_metrics(tune_res)

## # A tibble: 10 x 7
##   cost_complexity .metric .estimator  mean     n std_err .config
##   <dbl> <chr>   <chr>     <dbl> <int>   <dbl> <chr>
## 1 0.001   roc_auc hand_till 0.615     5 0.0184 Preprocessor1_Model01
## 2 0.00167  roc_auc hand_till 0.615     5 0.0184 Preprocessor1_Model02
## 3 0.00278  roc_auc hand_till 0.615     5 0.0184 Preprocessor1_Model03
## 4 0.00464  roc_auc hand_till 0.615     5 0.0201 Preprocessor1_Model04
## 5 0.00774  roc_auc hand_till 0.617     5 0.0184 Preprocessor1_Model05
## 6 0.0129   roc_auc hand_till 0.602     5 0.0226 Preprocessor1_Model06
## 7 0.0215   roc_auc hand_till 0.584     5 0.0204 Preprocessor1_Model07
## 8 0.0359   roc_auc hand_till 0.605     5 0.0132 Preprocessor1_Model08
## 9 0.0599   roc_auc hand_till 0.605     5 0.0126 Preprocessor1_Model09
## 10 0.1      roc_auc hand_till 0.5       5 0     Preprocessor1_Model10

arrange(tune_res)

## # Tuning results
## # 5-fold cross-validation using stratification
## # A tibble: 5 x 4
##   splits          id    .metrics      .notes
##   <list>         <chr> <list>        <list>
## 1 <split [289/75]> Fold1 <tibble [10 x 5]> <tibble [0 x 3]>
```

```

## 2 <split [291/73]> Fold2 <tibble [10 x 5]> <tibble [0 x 3]>
## 3 <split [291/73]> Fold3 <tibble [10 x 5]> <tibble [0 x 3]>
## 4 <split [292/72]> Fold4 <tibble [10 x 5]> <tibble [0 x 3]>
## 5 <split [293/71]> Fold5 <tibble [10 x 5]> <tibble [0 x 3]>


best_pruned <- select_best(tune_res)
class_tree_final <- finalize_workflow(class_tree_wf, best_pruned)
class_final_fit <- fit(class_tree_final, data = poke_train)
augment(class_final_fit, new_data = poke_test) %>% roc_auc(type_1, estimate = c(.pred_Bug, .pred_Fire,
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 roc_auc hand_till     0.628

auc = 0.6277206

class_final_fit %>%
  extract_fit_engine() %>%
  rpart.plot(roundint = FALSE)

```

## Warning: labs do not fit even at cex 0.15, there may be some overplotting



```

rf_spec <- rand_forest(mtry = tune()) %>%
  set_engine("ranger", importance = "impurity") %>%
  set_mode("classification")

rf_wf <- workflow() %>%
  add_model(rf_spec %>% set_args(trees = tune(), min_n = tune())) %>%
  add_recipe(pokeRecipe)

rf_grid <- grid_regular(mtry(range = c(1, 8)), trees(range = c(20, 100)), min_n(range = c(2, 40)), lev

```

mtry = number of predictors sampled at each split  
 trees = number of trees  
 min\_n = min number of observations required for another split


at mtry = 8 we are using all the predictors in each split. Therefore we can not go higher than 8, similarly, less than 1 is also not plausible because we won't have a useful model if 0 predictors are sampled at each split.

```

tune_res_rf <- tune_grid(
  rf_wf,
  resamples = poke_fold,
  grid = rf_grid,
  metrics = metric_set(roc_auc)
)

autoplot(tune_res_rf)

```



```

collect_metrics(tune_res_rf)

## # A tibble: 512 x 9
##   mtry trees min_n .metric .estimator  mean     n std_err .config
##   <int> <int> <int> <chr>   <chr>    <dbl> <int>   <dbl> <chr>
## 1     1     20      2 roc_auc hand_till  0.659     5 0.0203 Preprocessor1_Model~
## 2     2     20      2 roc_auc hand_till  0.673     5 0.0211 Preprocessor1_Model~
## 3     3     20      2 roc_auc hand_till  0.690     5 0.00448 Preprocessor1_Model~
## 4     4     20      2 roc_auc hand_till  0.689     5 0.0181 Preprocessor1_Model~
## 5     5     20      2 roc_auc hand_till  0.677     5 0.0202 Preprocessor1_Model~
## 6     6     20      2 roc_auc hand_till  0.676     5 0.0112 Preprocessor1_Model~
## 7     7     20      2 roc_auc hand_till  0.682     5 0.0163 Preprocessor1_Model~
## 8     8     20      2 roc_auc hand_till  0.695     5 0.0190 Preprocessor1_Model~
## 9     1     31      2 roc_auc hand_till  0.655     5 0.0170 Preprocessor1_Model~
## 10    2     31      2 roc_auc hand_till  0.679     5 0.0194 Preprocessor1_Model~
## # ... with 502 more rows

arrange(tune_res_rf)


## # Tuning results
## # 5-fold cross-validation using stratification
## # A tibble: 5 x 4
##   splits          id   .metrics       .notes
##   <list>         <chr> <list>        <list>
## 1 <split [289/75]> Fold1 <tibble [512 x 7]> <tibble [0 x 3]>
## 2 <split [291/73]> Fold2 <tibble [512 x 7]> <tibble [0 x 3]>
## 3 <split [291/73]> Fold3 <tibble [512 x 7]> <tibble [0 x 3]>
## 4 <split [292/72]> Fold4 <tibble [512 x 7]> <tibble [0 x 3]>
## 5 <split [293/71]> Fold5 <tibble [512 x 7]> <tibble [0 x 3]>

best_rf<- select_best(tune_res_rf)
rf_final <- finalize_workflow(rf_wf, best_rf)
rf_final_fit <- fit(rf_final, data = poke_train)
augment(rf_final_fit, new_data = poke_test) %>% roc_auc(type_1, estimate = c(.pred_Bug, .pred_Fire, .pred_Ghost))

## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>        <dbl>
## 1 roc_auc hand_till     0.678

rf_final_fit %>%
  extract_fit_parsnip() %>%
  vip()

```




special attack and attack were the most useful, generation was the least useful. This is expected since type heavily influences attack and special attack stats.

```
boost_spec <- boost_tree(trees = tune(), tree_depth = 4) %>%
  set_engine("xgboost") %>%
  set_mode("classification")

boost_wf <- workflow() %>%
  add_model(boost_spec %>% set_args(trees = tune())) %>%
  add_recipe(pokeRecipe)

boost_grid <- grid_regular(trees(range = c(10, 2000)), levels = 10)

boost_tune <- tune_grid(
  boost_wf,
  resamples = poke_fold,
  grid = boost_grid,
  metrics = metric_set(roc_auc)
)
autoplot(boost_tune)
```



```
collect_metrics(boost_tune)
```

```
## # A tibble: 10 x 7
##   trees .metric .estimator  mean     n std_err .config
##   <int> <chr>   <chr>    <dbl> <int>   <dbl> <chr>
## 1    10 roc_auc hand_till 0.682     5 0.0132 Preprocessor1_Model01
## 2   231 roc_auc hand_till 0.686     5 0.0157 Preprocessor1_Model02
## 3   452 roc_auc hand_till 0.688     5 0.0170 Preprocessor1_Model03
## 4   673 roc_auc hand_till 0.688     5 0.0179 Preprocessor1_Model04
## 5   894 roc_auc hand_till 0.687     5 0.0186 Preprocessor1_Model05
## 6  1115 roc_auc hand_till 0.687     5 0.0184 Preprocessor1_Model06
## 7  1336 roc_auc hand_till 0.687     5 0.0189 Preprocessor1_Model07
## 8  1557 roc_auc hand_till 0.687     5 0.0193 Preprocessor1_Model08
## 9  1778 roc_auc hand_till 0.686     5 0.0196 Preprocessor1_Model09
## 10 2000 roc_auc hand_till 0.686     5 0.0195 Preprocessor1_Model10
```

```
arrange(boost_tune)
```

```
## # Tuning results
## # 5-fold cross-validation using stratification
## # A tibble: 5 x 4
##   splits          id   .metrics      .notes
##   <list>         <chr> <list>       <list>
## 1 <split [289/75]> Fold1 <tibble [10 x 5]> <tibble [0 x 3]>
## 2 <split [291/73]> Fold2 <tibble [10 x 5]> <tibble [0 x 3]>
```

```
## 3 <split [291/73]> Fold3 <tibble [10 x 5]> <tibble [0 x 3]>
## 4 <split [292/72]> Fold4 <tibble [10 x 5]> <tibble [0 x 3]>
## 5 <split [293/71]> Fold5 <tibble [10 x 5]> <tibble [0 x 3]>

best_boost<- select_best(boost_tune)
boost_final <- finalize_workflow(boost_wf, best_boost)
boost_final_fit <- fit(boost_final, data = poke_train)
augment(boost_final_fit, new_data = poke_test) %>% roc_auc(type_1, estimate = c(.pred_Bug, .pred_Fire,
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 roc_auc hand_till     0.677
```

0.6771615